

Health AI – Intelligent Healthcare Assistant

1. Introduction:

- Project Title:Health Ai-intelligent healthcare assistant
- Team Member 1:Evangline.s
- Team Member 2:Maythini.R
- Team Member 3:Narmadha sri.k
- Team Member 4:Lavanya.k

2.Project overview:

Health AI is an intelligent healthcare assistant powered by IBM Granite models. It is designed to assist patients and healthcare providers by offering:

- Patient Chat: A conversational interface for patients to describe symptoms and get initial guidance.
- Disease Prediction: AI-driven predictions based on patient inputs and health data.
- Treatment Plan Suggestions: Providing treatment guidance and lifestyle recommendations.
- Extensibility: Developers can add new functionalities, like medicine reminders, medical history analysis, or integration with EHR systems.

The project leverages Google Colab for deployment, ensuring that even students or researchers without high-end machines can access GPU resources.By using IBM Granite LLMs from Hugging Face, the system combines language intelligence with healthcare applications, making medical information more accessible, understandable, and interactive.

3.Architecture

The architecture of Health AI follows a simple but efficient flow:

1. User Input Layer:

- Patients enter symptoms or queries via a Gradio interface.
- Input may include text like “I have fever, cough, and fatigue.”

2. Processing Layer:

- The system sends the input to an IBM Granite model hosted on Hugging Face.
- Pre-processing ensures text is cleaned and standardized.

3. Model Inference Layer:

- The Granite model generates a predicted response, which could be:
- Probable disease condition
- Suggested treatment approach
- Follow-up advice

4. Output Layer:

- Results are displayed back to the user in the Gradio interface.
- Users can interact in real-time, like chatting with a virtual doctor.

Tech Stack:

- Frontend: Gradio (for interactive UI)
- Backend: Python + Hugging Face (IBM Granite LLMs)
- Deployment: Google Colab (T4 GPU runtime)
- Version Control: GitHub

4.Setup Instructions:

- Pre-requisites
- Before running the project, ensure:
- Basic knowledge of Python and Gradio.
- Hugging Face account (for accessing IBM Granite models).
- Google Colab account (for running the notebook with GPU).
- GitHub account (for uploading project files).

Step-by-Step Setup

Step-1.Open Google Colab.

- Visit Google Colab.
- Create a new notebook.

Step 2: Change Runtime

- Go to Runtime → Change Runtime Type.
- Select GPU (T4) and save.

Step 3: Install Dependencies

- `!pip install transformers torch gradio -q`

Step 4: Load IBM Granite Model

- Login to Hugging Face.
- Select IBM Granite model (e.g., granite-3.2-2b-instruct).

Step 5: Run Application

- Execute the notebook cells.
- Gradio will generate a link → open in a new tab.

Step 6: Upload to GitHub

- Download notebook as .py or .ipynb.
- Push to GitHub repository.

5. Folder Structure

A well-defined folder structure ensures clarity, easy collaboration, and better project maintenance. The proposed structure for the HealthAI project is organized as follows:

- HealthAI/
The root directory of the project, containing all essential files and subfolders.
- Notebooks/
Contains Jupyter or Colab notebooks used for experimentation, prototyping, and testing.
- HealthAI.ipynb: The main notebook where initial model development, data exploration, or demonstrations are carried out.
- Src/This is the source code directory containing the main logic of the application.
- App.py: Defines the Gradio interface and application flow.
- Model.py: Handles model loading, initialization, and inference functions.
- Utils.py: Includes helper or utility functions to support preprocessing, evaluation, or other repetitive tasks.
- Models/Stores model-related files such as weights, checkpoints, or configuration details.
- Config.json: Contains model configuration parameters, hyperparameters, or setup details.
- Docs/Contains project-related documentation.
- ProjectReport.md: Extended documentation of the project, including design, architecture, and usage details.
- Screenshots/Stores visual assets such as interface previews, diagrams, or workflow images.
- Ui.png: Screenshot of the application's user interface.
- Requirements.txt
A list of all Python dependencies required to run the project (e.g., Gradio, Transformers, etc.).
- README.md

Provides a short overview of the project, setup instructions, and quick usage guidelines.

6. Running the Application

1. Open Google Colab Notebook → HealthAI.ipynb.
2. Install dependencies.
3. Load IBM Granite model from Hugging Face.
4. Start Gradio app → copy the generated URL.
5. Open in browser → interact with the chatbot.

7. API Documentation

Currently, the app is designed as an interactive tool, but can also act like an API.

- Endpoint 1 – Chat
- Input: Patient text (symptoms/query).
- Output: AI-generated conversational reply.

Endpoint 2 – Disease Prediction

- Input: Symptom description.
- Output: Possible diseases.

Endpoint 3 – Treatment Suggestion

- Input: Disease or symptom query.
- Output: Suggested treatments & lifestyle tips.

8. Authentication

- Hugging Face API Key required to download IBM Granite models.
- Store API keys securely (never upload to GitHub).

Use environment variables in Colab:

```
Import os os.environ["HF_TOKEN"] = "your_huggingface_token"
```

9. User Interface

Built with Gradio, which provides a simple UI:

- Text box for inputting patient symptoms.

- Chat window for conversational interaction.
- Results panel for displaying predictions and treatment plans.
- Mobile-friendly and lightweight.

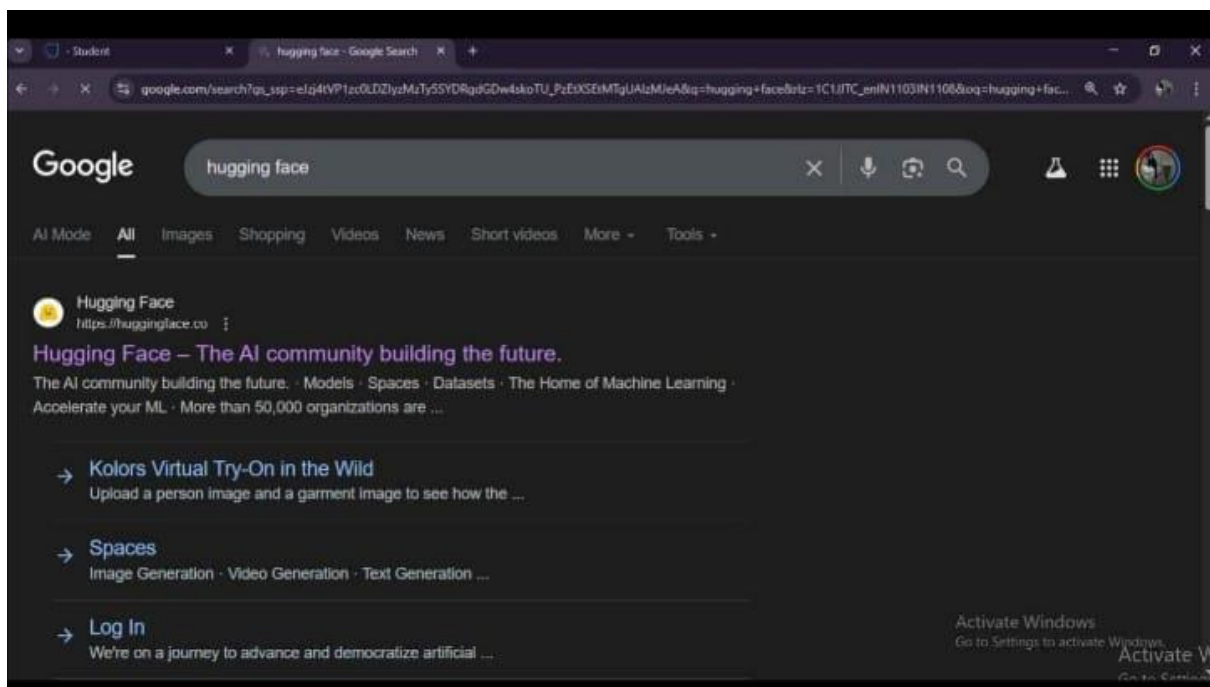
10. Testing

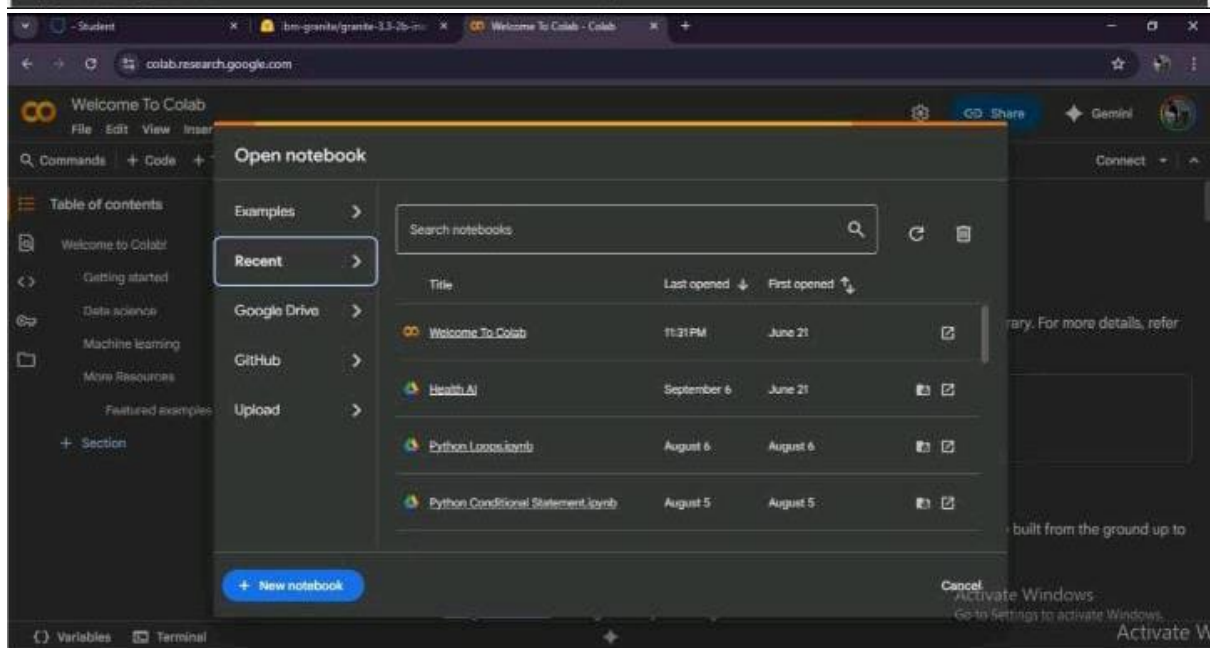
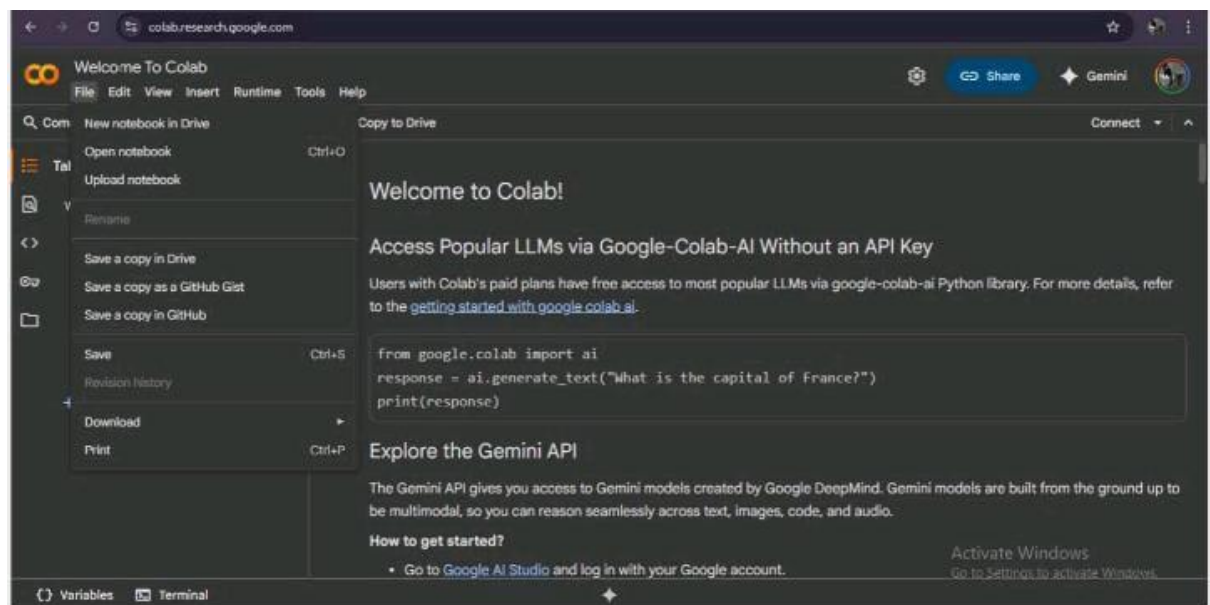
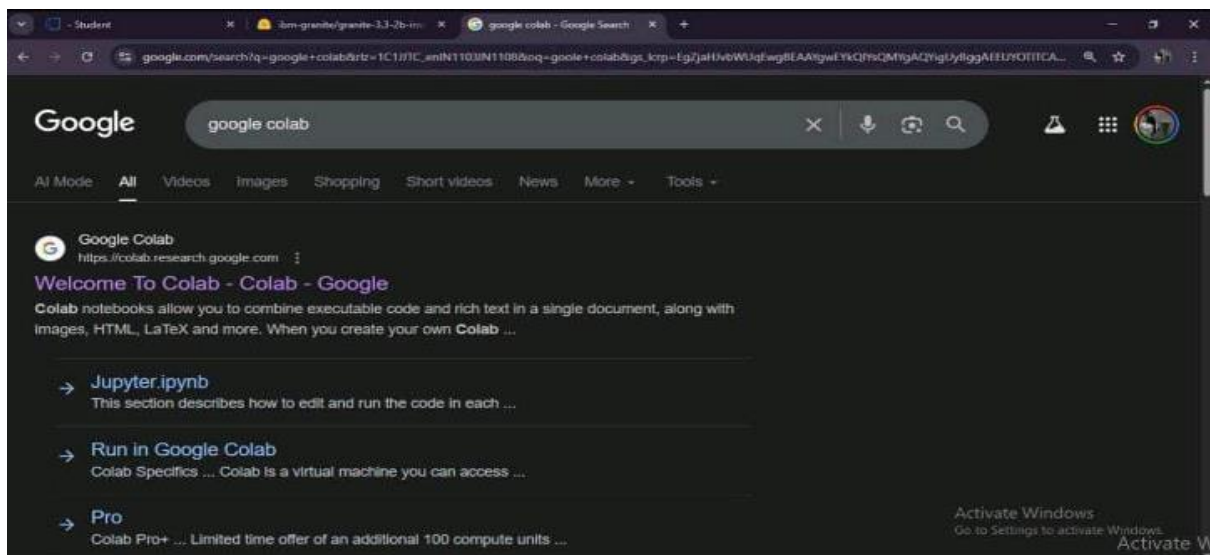
- Unit Testing: Validate each function in utils.py.
- Integration Testing: Ensure Hugging Face API works with Colab.
- User Testing: Try queries like:

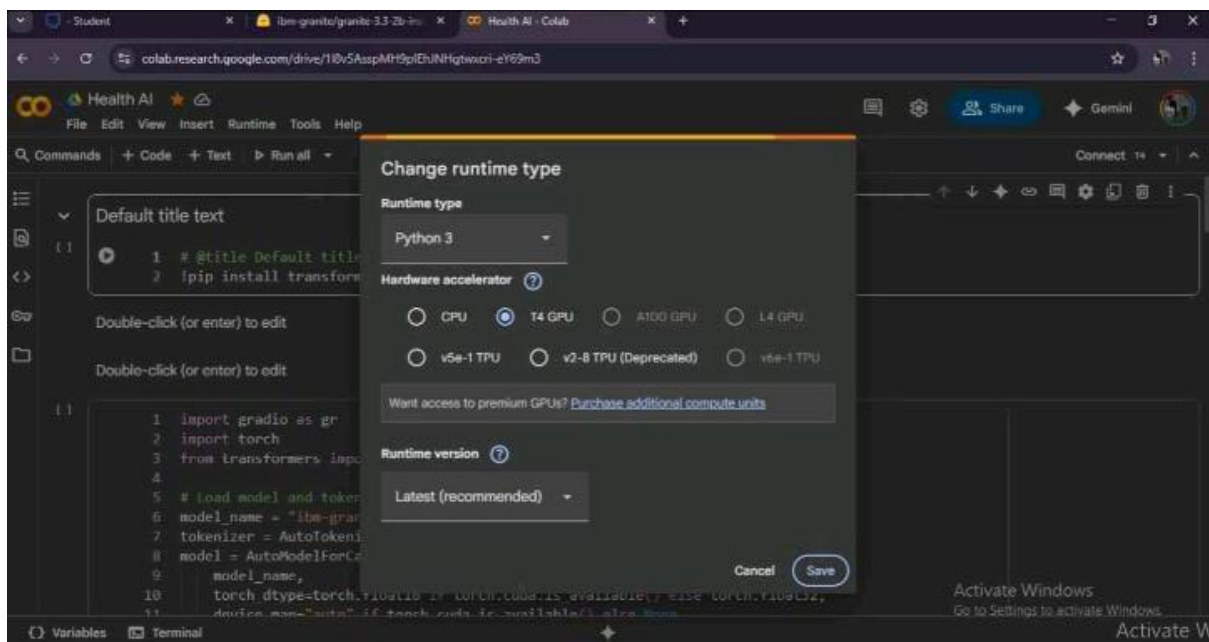
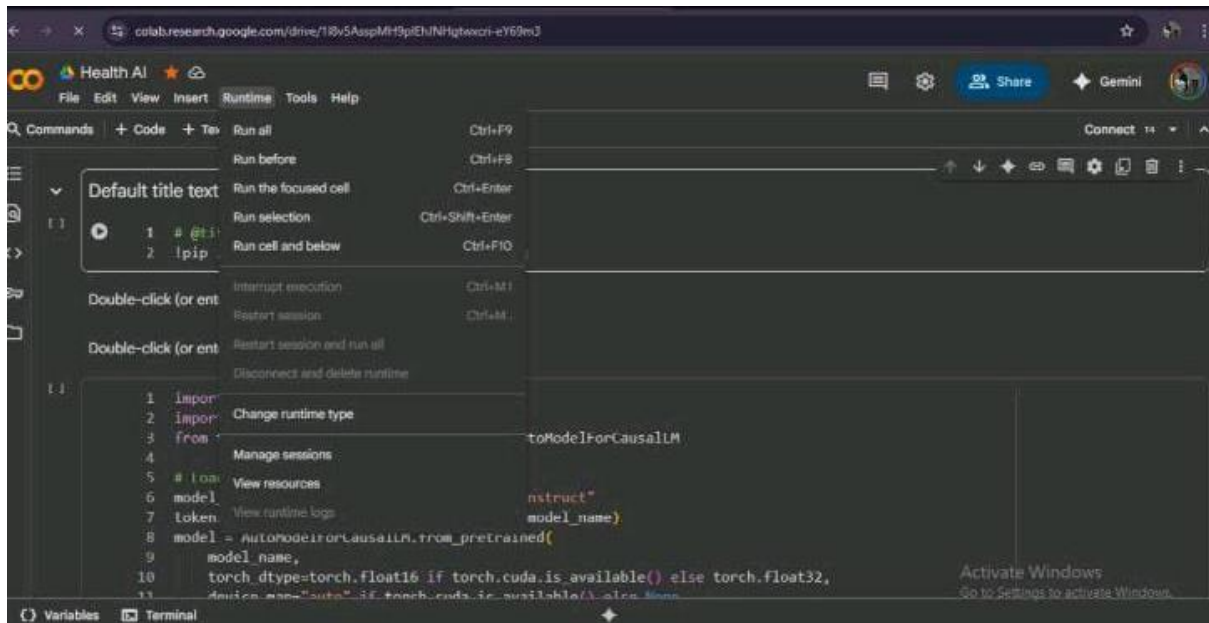
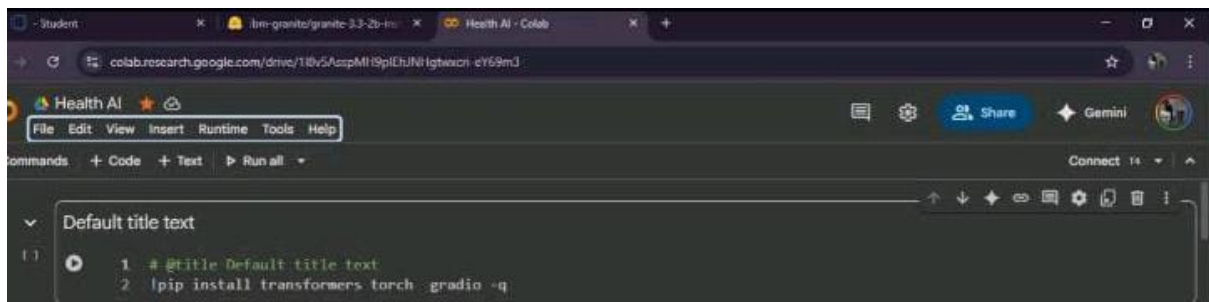
“I have cough and fever” → Predict flu or cold.

“What is the treatment for diabetes?” → Suggest lifestyle + medication advice.

11. Screenshots/demo








```
colab.research.google.com/drive/18v5AaspMH9pBtUHNHgtwscii-eY63m3
Health AI
File Edit View Insert Runtime Tools Help
Q Commands + Code + Text ▶ Run all
Connect 11
Double-click (or enter) to edit
11
1 import gradio as gr
2 import torch
3 from transformers import AutoTokenizer, AutoModelForCausalLM
4
5 # load model and tokenizer
6 model_name = "ibm-granite/granite-3.2-2b-instruct"
7 tokenizer = AutoTokenizer.from_pretrained(model_name)
8 model = AutoModelForCausalLM.from_pretrained(
9     model_name,
10     torch_dtype=torch.float16 if torch.cuda.is_available() else torch.float32,
11     device_map="auto" if torch.cuda.is_available() else None
12 )
13
14 if tokenizer.pad_token is None:
15     tokenizer.pad_token = tokenizer.eos_token
16
17 def generate_response(prompt, max_length=1024):
18     inputs = tokenizer(prompt, return_tensors="pt", truncation=True, max_length=512)
```

```
Health AI
File Edit View Insert Runtime Tools Help
Q Commands + Code + Text ▶ Run all
Connect 14
16
17 def generate_response(prompt, max_length=1024):
18     inputs = tokenizer(prompt, return_tensors="pt", truncation=True, max_length=512)
19
20     if torch.cuda.is_available():
21         inputs = {k: v.to(model.device) for k, v in inputs.items()}
22
23     with torch.no_grad():
24         outputs = model.generate(
25             **inputs,
26             max_length=max_length,
27             temperature=0.7,
28             do_sample=True,
29             pad_token_id=tokenizer.eos_token_id
30         )
31
32     response = tokenizer.decode(outputs[0], skip_special_tokens=True)
33     response = response.replace(prompt, "").strip()
34     return response
35
36 def disease_prediction(symptoms):
37     prompt = f"Based on the following symptoms, provide possible medical conditions and general medication suggestions. Always emphasize the importance of consulting a doctor for proper diagnosis and treatment. This is for informational purposes only. Please consult a healthcare professional for proper diagnosis and treatment."
38     return generate_response(prompt, max_length=1200)
```

```
colab.research.google.com/drive/18v5AaspMH9pBtUHNHgtwscii-eY63m3
Health AI
File Edit View Insert Runtime Tools Help
Q Commands + Code + Text ▶ Run all
Connect 14
34 return response
35
36 def disease_prediction(symptoms):
37     prompt = f"Based on the following symptoms, provide possible medical conditions and general medication suggestions. Always emphasize the importance of consulting a doctor for proper diagnosis and treatment. This is for informational purposes only. Please consult a healthcare professional for proper diagnosis and treatment."
38     return generate_response(prompt, max_length=1200)
39
40 def treatment_plan(condition, age, gender, medical_history):
41     prompt = f"Generate personalized treatment suggestions for the following patient information. Include home remedies and general medication guidelines.\n\nMedical Condition: {condition}\n\nAge: {age}\n\nGender: {gender}\n\nMedical History: {medical_history}\n\nPersonalized treatment plan including home remedies and medication guidelines:\n\n**IMPORTANT: This is for informational purposes only. Please consult a healthcare professional for proper treatment.**\n\nTreatment Plan:"
42     return generate_response(prompt, max_length=1200)
43
44 # Create Gradio Interface
45 with gr.Blocks() as app:
46     gr.Markdown("# Medical AI Assistant")
47     gr.Markdown("**Disclaimer: This is for informational purposes only. Always consult healthcare
```



```
colab.research.google.com/drive/1Bv5AaspMr9pEh1H4Hgtw01-eY69m3

Health AI
File Edit View Insert Runtime Tools Help

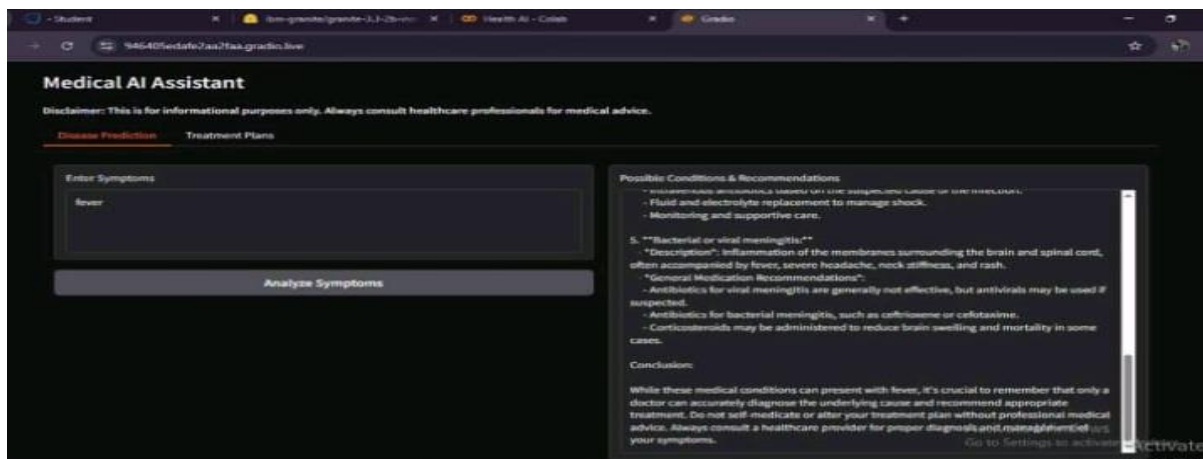
Commands + Code + Text Run all Connect 14

11 48 professionals for medical advice.**")
49
50 with gr.Tabs():
51     with gr.Tabitem("Disease Prediction"):
52         with gr.Column():
53             symptoms_input = gr.Textbox(
54                 label="Enter Symptoms",
55                 placeholder="e.g., fever, headache, cough, fatigue...",
56                 lines=4
57             )
58             predict_btn = gr.Button("Analyze Symptoms")
59
60         with gr.Column():
61             prediction_output = gr.Textbox(label="Possible Conditions & Recommendations",
62                 lines=20)
63
64             predict_btn.click(disease_prediction, inputs=symptoms_input, outputs=prediction_output)
65
66     with gr.Tabitem("Treatment Plans"):
67         with gr.Column():
68             condition_input = gr.Textbox(
69                 label="Medical Condition",
70                 placeholder="e.g., diabetes, hypertension, migraine...",
71                 lines=2
72             )
73             age_input = gr.Number(label="Age", value=30)
74             gender_input = gr.Dropdown(
75                 choices=["Male", "Female", "Other"],
76                 label="Gender",
77                 value="Male"
78             )
79             history_input = gr.Textbox(
80                 label="Medical History",
81                 placeholder="Previous conditions, allergies, medications or None",
82                 lines=3
83             )
84             plan_btn = gr.Button("Generate Treatment Plan")
85
86         with gr.Column():
87             plan_output = gr.Textbox(label="Personalized Treatment Plan", lines=20)
88
89             plan_btn.click(treatment_plan, inputs=[condition_input, age_input, gender_input,
90                 history_input], outputs=plan_output)
```

```
Health AI
File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Connect 14

69
70 label="Medical Condition",
71 placeholder="e.g., diabetes, hypertension, migraine...",
72 lines=2
73 )
74 age_input = gr.Number(label="Age", value=30)
75 gender_input = gr.Dropdown(
76     choices=["Male", "Female", "Other"],
77     label="Gender",
78     value="Male"
79 )
80 history_input = gr.Textbox(
81     label="Medical History",
82     placeholder="Previous conditions, allergies, medications or None",
83     lines=3
84 )
85 plan_btn = gr.Button("Generate Treatment Plan")
86
87 with gr.Column():
88     plan_output = gr.Textbox(label="Personalized Treatment Plan", lines=20)
89
90 plan_btn.click(treatment_plan, inputs=[condition_input, age_input, gender_input,
91     history_input], outputs=plan_output)
```



12. Known Issues

- Dependency on Google Colab GPU (sessions may timeout).
- Hugging Face model load time is slow initially.
- Limited to text-only input (no speech/image support yet).
- Predictions are not medical diagnoses – only guidance.

13. Future Enhancements

- Deploy on IBM Cloud for better scalability.
- Voice-based assistant for accessibility.
- Integration with wearables (smartwatch, fitness tracker).
- Multilingual support for rural/remote users.
- Larger medical dataset for better disease coverage.